

Task2

September 28, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
import plotly.graph_objects as go
import plotly.offline as pyo
pyo.init_notebook_mode()
import duckdb
import warnings
warnings.filterwarnings("ignore")
```

1 Checking Data

```
[2]: transactions = pd.read_csv('transactions_1.csv')
transactions
```

```
[2]:
```

	time	status	f0_
0	00h 00	approved	9
1	00h 00	denied	6
2	00h 00	refunded	1
3	00h 01	denied	8
4	00h 01	approved	13
...
4229	23h 59	processing	4
4230	23h 59	denied	1
4231	23h 59	backend_reversed	5
4232	23h 59	approved	10
4233	23h 59	reversed	6

[4234 rows x 3 columns]

```
[3]: transactions2 = pd.read_csv('transactions_2.csv')
transactions2
```

```
[3]:
```

	time	status	count
0	00h 00	reversed	7
1	00h 00	approved	9

2	00h 00	processing	12
3	00h 00	denied	3
4	00h 00	backend_reversed	2
...
3939	23h 57	denied	11
3940	23h 58	denied	4
3941	23h 58	approved	35
3942	23h 59	denied	10
3943	23h 59	approved	38

[3944 rows x 3 columns]

2 Checking for Problems in Data

```
[4]: print('_ % NaNs __ First File _____')
for col in transactions.columns:
    pct = transactions[f'{col}'].isnull().sum() * 100 / len(transactions.index)
    print(f'{col} - {pct}%')
print('\n_ % NaNs __ Second File _____')
for col in transactions2.columns:
    pct = transactions2[f'{col}'].isnull().sum() * 100 / len(transactions2.
    ↪index)
    print(f'{col} - {pct}%')
```

```
_ % NaNs __ First File _____
time - 0.0%
status - 0.0%
f0_ - 0.0%

_ % NaNs __ Second File _____
time - 0.0%
status - 0.0%
count - 0.0%
```

```
[5]: print('_ % Zeros __ First File _____')
for col in transactions.columns:
    pct = len(transactions[transactions[f'{col}']==0].index) * 100 /
    ↪len(transactions.index)
    print(f'{col} - {pct}%')
print('\n_ % Zeros __ Second File _____')
for col in transactions2.columns:
    pct = len(transactions2[transactions2[f'{col}']==0].index) * 100 /
    ↪len(transactions2.index)
    print(f'{col} - {pct}%')
```

```
_ % Zeros __ First File _____
time - 0.0%
```

```

status - 0.0%
f0_ - 0.0%

_ % Zeros __ Second File -----
time - 0.0%
status - 0.0%
count - 0.0%

```

```

[6]: print('_ % Duplicates Row Wise __ First File -----')
      for col in transactions.columns:
          non_duplicates = len(transactions.drop_duplicates(keep='first').index)
          duplicates = len(transactions.index) - non_duplicates
          pct = duplicates * 100 / len(transactions.index)
          print(f'{col} - {pct}%')

      print('\n_ % Duplicates Row Wise __ Second File -----')
      for col in transactions2.columns:
          non_duplicates = len(transactions2.drop_duplicates(keep='first').index)
          duplicates = len(transactions2.index) - non_duplicates
          pct = duplicates * 100 / len(transactions2.index)
          print(f'{col} - {pct}%')

```

```

_ % Duplicates Row Wise __ First File -----
time - 0.0%
status - 0.0%
f0_ - 0.0%

```

```

_ % Duplicates Row Wise __ Second File -----
time - 0.0%
status - 0.0%
count - 0.0%

```

```

[7]: print('_ % Duplicates Column Wise __ First File -----')
      for col in transactions.columns:
          non_duplicates = len(transactions.drop_duplicates(col, keep='first').index)
          duplicates = len(transactions.index) - non_duplicates
          pct = duplicates * 100 / len(transactions.index)
          print(f'{col} - {pct}%')

      print('\n_ % Duplicates Column Wise __ Second File -----')
      for col in transactions2.columns:
          non_duplicates = len(transactions2.drop_duplicates(col, keep='first').index)
          duplicates = len(transactions2.index) - non_duplicates
          pct = duplicates * 100 / len(transactions2.index)
          print(f'{col} - {pct}%')

```

```

_ % Duplicates Column Wise __ First File -----
time - 68.77657061880019%
status - 99.83467170524327%

```

```
f0_ - 90.3873405762872%
```

```
_ % Duplicates Column Wise __ Second File _  
time - 66.55679513184585%  
status - 99.82251521298174%  
count - 88.56490872210954%
```

3 Analysing Data Similarity

3.1 Transforming and preparing data

```
[8]: transactions.rename(columns={'f0_': 'count'}, inplace=True)
```

```
[9]: transactions = transactions.pivot_table(['count'], 'time', 'status').  
    ↪droplevel(0,axis=1)  
transactions.columns =  
    ↪['approved', 'backend_reversed', 'denied', 'failed', 'processing', 'refunded', 'reversed']  
transactions.fillna(0, inplace=True)
```

```
[10]: transactions2 = transactions2.pivot_table(['count'], 'time', 'status').  
    ↪droplevel(0,axis=1)  
transactions2.columns =  
    ↪['approved', 'backend_reversed', 'denied', 'failed', 'processing', 'refunded', 'reversed']  
transactions2.fillna(0, inplace=True)
```

```
[11]: transactions
```

```
[11]:
```

	approved	backend_reversed	denied	failed	processing	refunded	\
time							
00h 00	9.0	0.0	6.0	0.0	0.0	1.0	
00h 01	13.0	0.0	8.0	0.0	0.0	0.0	
00h 02	11.0	0.0	7.0	0.0	0.0	1.0	
00h 03	12.0	0.0	3.0	0.0	0.0	0.0	
00h 04	11.0	0.0	0.0	0.0	0.0	0.0	
...		
23h 55	32.0	0.0	4.0	0.0	0.0	0.0	
23h 56	30.0	0.0	3.0	0.0	0.0	0.0	
23h 57	21.0	0.0	4.0	0.0	0.0	0.0	
23h 58	11.0	4.0	3.0	0.0	3.0	1.0	
23h 59	10.0	5.0	1.0	0.0	4.0	0.0	

	reversed
time	
00h 00	0.0
00h 01	0.0
00h 02	0.0
00h 03	0.0

```

00h 04      0.0
...
23h 55      0.0
23h 56      0.0
23h 57      1.0
23h 58      2.0
23h 59      6.0

```

[1322 rows x 7 columns]

```
[12]: transactions.iloc[800:1000]
```

```

[12]:      approved  backend_reversed  denied  failed  processing  refunded  \
time
14h 40      17.0           4.0      2.0      0.0           8.0          0.0
14h 41     248.0           1.0     62.0      0.0           7.0          2.0
14h 42     434.0           3.0     87.0      0.0           0.0         13.0
14h 43     782.0           7.0     98.0      0.0           0.0         11.0
14h 44     305.0           6.0     54.0      1.0          14.0          1.0
...
18h 20     227.0          14.0     40.0      0.0           3.0          4.0
18h 22      49.0           5.0     14.0      0.0           3.0          2.0
18h 23     100.0          10.0     16.0      0.0           3.0          3.0
18h 24     306.0          11.0     60.0      0.0           0.0          5.0
18h 25      28.0           4.0      5.0      0.0           1.0          1.0

```

```

      reversed
time
14h 40      31.0
14h 41     180.0
14h 42     193.0
14h 43     109.0
14h 44      82.0
...
18h 20     101.0
18h 22      18.0
18h 23      35.0
18h 24     103.0
18h 25       7.0

```

[200 rows x 7 columns]

3.2 Kolmogorov–Smirnov test

```
[13]: from scipy.stats import ks_2samp

for status in transactions.columns:
    first_sample = transactions[status]
    second_sample = transactions2[status]

    print(status, '\t', ks_2samp(first_sample, second_sample)[1], '\t\t\t',
          ↪ks_2samp(first_sample, second_sample)[1]>0.05)
```

approved	2.749048125448754e-13	False	
backend_reversed	2.8130173420842248e-27		False
denied	6.314987236267669e-07	False	
failed	1.3325827491135293e-07	False	
processing	8.998373686853123e-10	False	
refunded	0.04250709469575456	False	
reversed	5.186332038015167e-12	False	

3.3 Distributions per Status

```
[14]: fig = plt.figure(figsize=(32,32))
fig.subplots_adjust(hspace=0.2, wspace=0.2)

for i, status in enumerate(transactions.columns):


    ax = fig.add_subplot(3, 3, i+1)

    first_sample = transactions[status]
    second_sample = transactions2[status]

    sns.distplot(x=first_sample, kde=True, bins=30, ax=ax, color='red')
    ax.axvline(np.mean(first_sample) + np.std(first_sample), color='red')
    sns.distplot(x=second_sample, kde=True, bins=30, ax=ax, color='blue')
    ax.axvline(np.mean(second_sample) + np.std(second_sample), color='blue')

    ax.set_title(str(status))


plt.show()
```



3.4 Correlation


```
[15]: fig = plt.figure(figsize=(20,20))
corr_df = pd.merge(transactions, transactions2, how='left', on='time',
    ↳ suffixes=['_1', '_2'])
sns.heatmap(corr_df.corr(), annot=True, fmt='.1f')
plt.title('Correlation Map')
```

```
[15]: Text(0.5, 1.0, 'Correlation Map')
```



```
[16]: fig = plt.figure(figsize=(20,20))
corr_df = pd.merge(transactions, transactions2, how='left', on='time',
    ↳ suffixes=['_1', '_2'])
sns.heatmap(corr_df.cov(), annot=True, fmt='.1f', vmin=0, vmax=5*10**3)
plt.title('Covariance Map')
```

```
[16]: Text(0.5, 1.0, 'Covariance Map')
```

Considering our context with **the first task**, and what we have analyzed, I think it is fair to assume these are from the same context and POS

Reasons to believe so:

- High correlations between approved 1 and 2, denied 1 and approved 2... basically high correlations between main data features
- High covariance between Approved sales
- With the context that other status are failures or anomalies, it makes sense that there will be divergences between these features
- Similar distributions even though the Kolmogorov test failed, with variances probably being normal and/or explainable

4 Data Augmentation

We will be adding various feature engineering columns to better pronounce the anomalies in the data, but first we will do some necessary basic transformations.

```
[17]: transactions = pd.concat([transactions.sort_index(ascending=True),  
    ↪ transactions2.sort_index(ascending=True)])  
duckdb.query('''  
SELECT * FROM transactions  
''')
```

[17]:

	approved reversed double double	backend_reversed double	denied double	failed double	processing double	refunded double
	9.0	0.0	6.0	0.0	0.0	1.0
0.0	13.0	0.0	8.0	0.0	0.0	0.0
0.0	11.0	0.0	7.0	0.0	0.0	1.0
0.0	12.0	0.0	3.0	0.0	0.0	0.0
0.0	11.0	0.0	0.0	0.0	0.0	0.0
0.0	9.0	0.0	2.0	0.0	0.0	0.0
0.0	9.0	0.0	2.0	0.0	0.0	0.0
0.0	11.0	0.0	1.0	0.0	0.0	0.0
1.0	7.0	0.0	2.0	0.0	0.0	0.0
0.0	9.0	0.0	0.0	0.0	0.0	0.0
0.0
.
.
.	33.0	0.0	11.0	0.0	0.0	0.0
1.0	41.0	0.0	8.0	0.0	0.0	0.0

0.0						
	43.0	0.0	9.0	0.0	0.0	0.0
0.0						
	33.0	0.0	7.0	0.0	0.0	1.0
0.0						
	39.0	0.0	4.0	0.0	0.0	1.0
0.0						
	31.0	0.0	8.0	0.0	0.0	0.0
0.0						
	34.0	0.0	8.0	0.0	0.0	0.0
0.0						
	40.0	0.0	11.0	0.0	0.0	0.0
0.0						
	35.0	0.0	4.0	0.0	0.0	0.0
0.0						
	38.0	0.0	10.0	0.0	0.0	0.0
0.0						

2641 rows (20 shown)
columns

7

```
[18]: transactions.index = pd.to_datetime(transactions.index, format='%Hh %M').
      ↪strftime('%H:%M')
```

```
[19]: transactions = transactions.reset_index()
      transactions['hour'] = pd.to_datetime(transactions.time).dt.hour
```

4.0.1 Creating boolean columns for each hour

This is a demonstration, it will be added after we have added and evaluated our other new features

```
[20]: pd.get_dummies(transactions, columns=['hour'])
```

```
[20]:
```

	time	approved	backend_reversed	denied	failed	processing	refunded	\
0	00:00	9.0	0.0	6.0	0.0	0.0	1.0	
1	00:01	13.0	0.0	8.0	0.0	0.0	0.0	
2	00:02	11.0	0.0	7.0	0.0	0.0	1.0	
3	00:03	12.0	0.0	3.0	0.0	0.0	0.0	
4	00:04	11.0	0.0	0.0	0.0	0.0	0.0	
...	
2636	23:55	31.0	0.0	8.0	0.0	0.0	0.0	
2637	23:56	34.0	0.0	8.0	0.0	0.0	0.0	
2638	23:57	40.0	0.0	11.0	0.0	0.0	0.0	
2639	23:58	35.0	0.0	4.0	0.0	0.0	0.0	
2640	23:59	38.0	0.0	10.0	0.0	0.0	0.0	

	reversed	hour_0	hour_1	...	hour_14	hour_15	hour_16	hour_17	\
0	0.0	True	False	...	False	False	False	False	
1	0.0	True	False	...	False	False	False	False	
2	0.0	True	False	...	False	False	False	False	
3	0.0	True	False	...	False	False	False	False	
4	0.0	True	False	...	False	False	False	False	
...	
2636	0.0	False	False	...	False	False	False	False	
2637	0.0	False	False	...	False	False	False	False	
2638	0.0	False	False	...	False	False	False	False	
2639	0.0	False	False	...	False	False	False	False	
2640	0.0	False	False	...	False	False	False	False	

	hour_18	hour_19	hour_20	hour_21	hour_22	hour_23
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
...
2636	False	False	False	False	False	True
2637	False	False	False	False	False	True
2638	False	False	False	False	False	True
2639	False	False	False	False	False	True
2640	False	False	False	False	False	True

[2641 rows x 32 columns]

Denied and *Approved* transactions are very strongly correlated. With this, we can create a Linear Regression model to predict the denied transactions based on the approved ones. If the model makes a prediction of denials, and the reality is too far from the prediction, it should be considered anomalous. The distance will be calculated by the square difference between prediction and reality

```
[21]: from sklearn.linear_model import LinearRegression

epsilon = 0.001 # number to avoid infinity in divisions by 0

roll_window = 15 # window for rolling calculations

transactions['total'] = (transactions.approved
                        + transactions.backend_reversed
                        + transactions.denied
                        + transactions.failed
                        + transactions.processing
                        + transactions.refunded
                        + transactions.reversed
```

```

    )

transactions['denied_approved'] = (transactions.denied / (transactions.
    ↪approved+epsilon))*100
transactions['denied_pct'] = (transactions.denied / (transactions.total))*100

transactions['reversed_approved'] = (transactions.reversed / (transactions.
    ↪approved+epsilon))*100
transactions['reversed_pct'] = (transactions.reversed / (transactions.
    ↪total))*100

transactions['failed_approved'] = (transactions.failed / (transactions.
    ↪approved+epsilon))*100
transactions['failed_pct'] = (transactions.failed / (transactions.total))*100

transactions['processing_approved'] = (transactions.processing / (transactions.
    ↪approved+epsilon))*100
transactions['processing_pct'] = (transactions.processing / (transactions.
    ↪total))*100

transactions['refunded_approved'] = (transactions.refunded / (transactions.
    ↪approved+epsilon))*100
transactions['refunded_pct'] = (transactions.refunded / (transactions.
    ↪total))*100

transactions['backend_reversed_approved'] = (transactions.backend_reversed /
    ↪(transactions.approved+epsilon))*100
transactions['backend_reversed_pct'] = (transactions.backend_reversed /
    ↪(transactions.total))*100

transactions['denied_sqr'] = transactions.denied**2
transactions['reversed_sqr'] = transactions.reversed**2
transactions['failed_sqr'] = transactions.failed**2
transactions['processing_sqr'] = transactions.processing**2
transactions['refunded_sqr'] = transactions.refunded**2
transactions['backend_reversed_sqr'] = transactions.backend_reversed**2

transactions['denied_sqroot'] = transactions.denied**0.5
transactions['reversed_sqroot'] = transactions.reversed**0.5
transactions['failed_sqroot'] = transactions.failed**0.5
transactions['processing_sqroot'] = transactions.processing**0.5
transactions['refunded_sqroot'] = transactions.refunded**0.5
transactions['backend_reversed_sqroot'] = transactions.backend_reversed**0.5

```

```

transactions['denied_approved_roll30_corr'] = transactions.denied.
    ↳rolling(roll_window).corr(transactions.approved)
transactions['reversed_approved_roll30_corr'] = transactions.reversed.
    ↳rolling(roll_window).corr(transactions.approved)
transactions['failed_approved_roll30_corr'] = transactions.failed.
    ↳rolling(roll_window).corr(transactions.approved)
transactions['processing_approved_roll30_corr'] = transactions.processing.
    ↳rolling(roll_window).corr(transactions.approved)
transactions['refunded_approved_roll30_corr'] = transactions.refunded.
    ↳rolling(roll_window).corr(transactions.approved)
transactions['backend_reversed_approved_roll30_corr'] = transactions.
    ↳backend_reversed.rolling(roll_window).corr(transactions.approved)

transactions['denied_roll30_cumsum'] = transactions.denied.rolling(roll_window).
    ↳sum()
transactions['reversed_roll30_cumsum'] = transactions.reversed.
    ↳rolling(roll_window).sum()
transactions['failed_roll30_cumsum'] = transactions.failed.rolling(roll_window).
    ↳sum()
transactions['processing_roll30_cumsum'] = transactions.processing.
    ↳rolling(roll_window).sum()
transactions['refunded_roll30_cumsum'] = transactions.refunded.
    ↳rolling(roll_window).sum()
transactions['backend_reversed_roll30_cumsum'] = transactions.backend_reversed.
    ↳rolling(roll_window).sum()

transactions['denied_roll30_avg'] = transactions.denied.rolling(roll_window).
    ↳mean()
transactions['reversed_roll30_avg'] = transactions.reversed.
    ↳rolling(roll_window).mean()
transactions['failed_roll30_avg'] = transactions.failed.rolling(roll_window).
    ↳mean()
transactions['processing_roll30_avg'] = transactions.processing.
    ↳rolling(roll_window).mean()
transactions['refunded_roll30_avg'] = transactions.refunded.
    ↳rolling(roll_window).mean()
transactions['backend_reversed_roll30_avg'] = transactions.backend_reversed.
    ↳rolling(roll_window).mean()

transactions['denied_roll30_min'] = transactions.denied.rolling(roll_window).
    ↳min()
transactions['reversed_roll30_min'] = transactions.reversed.
    ↳rolling(roll_window).min()

```

```

transactions['failed_roll30_min'] = transactions.failed.rolling(roll_window).
    ↳mean()
transactions['processing_roll30_min'] = transactions.processing.
    ↳rolling(roll_window).min()
transactions['refunded_roll30_min'] = transactions.refunded.
    ↳rolling(roll_window).min()
transactions['backend_reversed_roll30_min'] = transactions.backend_reversed.
    ↳rolling(roll_window).min()

transactions['denied_roll30_max'] = transactions.denied.rolling(roll_window).
    ↳max()
transactions['reversed_roll30_max'] = transactions.reversed.
    ↳rolling(roll_window).max()
transactions['failed_roll30_max'] = transactions.failed.rolling(roll_window).
    ↳max()
transactions['processing_roll30_max'] = transactions.processing.
    ↳rolling(roll_window).max()
transactions['refunded_roll30_max'] = transactions.refunded.
    ↳rolling(roll_window).max()
transactions['backend_reversed_roll30_max'] = transactions.backend_reversed.
    ↳rolling(roll_window).max()

# For denied transactions, we will use linear regression to measure distance of
↳data point from hypothesis
X, y = np.array(transactions.approved).reshape(-1, 1), transactions.denied
linearR_denied = LinearRegression().fit(X,y)
score = linearR_denied.score(X,y)
print(f'Linear Regression score: {score}')
predicted_denied = linearR_denied.predict(X)

transactions['denied_sqr_distance'] = (transactions.denied -
    ↳predicted_denied)**2

transactions = transactions.replace(np.nan, 0)

```

Linear Regression score: 0.7303313244160479

4.0.2 Checking for NaNs post feature engineering

```

[22]: print('\t\t\t___NaNs___\n')
for col in transactions.columns:
    pct = transactions[f'{col}'].isnull().sum() * 100 / len(transactions.index)
    print(f'{col} - {pct}%')

```

___NaNs___

time - 0.0%
approved - 0.0%
backend_reversed - 0.0%
denied - 0.0%
failed - 0.0%
processing - 0.0%
refunded - 0.0%
reversed - 0.0%
hour - 0.0%
total - 0.0%
denied_approved - 0.0%
denied_pct - 0.0%
reversed_approved - 0.0%
reversed_pct - 0.0%
failed_approved - 0.0%
failed_pct - 0.0%
processing_approved - 0.0%
processing_pct - 0.0%
refunded_approved - 0.0%
refunded_pct - 0.0%
backend_reversed_approved - 0.0%
backend_reversed_pct - 0.0%
denied_sqr - 0.0%
reversed_sqr - 0.0%
failed_sqr - 0.0%
processing_sqr - 0.0%
refunded_sqr - 0.0%
backend_reversed_sqr - 0.0%
denied_sqroot - 0.0%
reversed_sqroot - 0.0%
failed_sqroot - 0.0%
processing_sqroot - 0.0%
refunded_sqroot - 0.0%
backend_reversed_sqroot - 0.0%
denied_approved_roll30_corr - 0.0%
reversed_approved_roll30_corr - 0.0%
failed_approved_roll30_corr - 0.0%
processing_approved_roll30_corr - 0.0%
refunded_approved_roll30_corr - 0.0%
backend_reversed_approved_roll30_corr - 0.0%
denied_roll30_cumsum - 0.0%
reversed_roll30_cumsum - 0.0%
failed_roll30_cumsum - 0.0%
processing_roll30_cumsum - 0.0%
refunded_roll30_cumsum - 0.0%
backend_reversed_roll30_cumsum - 0.0%
denied_roll30_avg - 0.0%
reversed_roll30_avg - 0.0%


```
failed_roll30_avg - 0.0%
processing_roll30_avg - 0.0%
refunded_roll30_avg - 0.0%
backend_reversed_roll30_avg - 0.0%
denied_roll30_min - 0.0%
reversed_roll30_min - 0.0%
failed_roll30_min - 0.0%
processing_roll30_min - 0.0%
refunded_roll30_min - 0.0%
backend_reversed_roll30_min - 0.0%
denied_roll30_max - 0.0%
reversed_roll30_max - 0.0%
failed_roll30_max - 0.0%
processing_roll30_max - 0.0%
refunded_roll30_max - 0.0%
backend_reversed_roll30_max - 0.0%
denied_sqr_distance - 0.0%
```

4.0.3 Distribution of variables

```
[23]: fig = plt.figure(figsize=(32,32))
fig.subplots_adjust(hspace=0.5, wspace=0.2)
for i, col in enumerate(transactions.columns):
    if col=='time':continue

    ax = fig.add_subplot(9, 9, i+1)
    sns.distplot(transactions[col], kde=False, ax=ax)

    ax.set_title(col)

plt.show()
```



```
2640    23
Name: hour, Length: 2641, dtype: object
```

```
[25]: transactions.drop('time', axis=1, inplace=True)
```

One Hot Encoder for each Hour

```
[26]: transactions = pd.get_dummies(transactions, 'hour', dtype=int)
transactions
```

```
[26]:
```

	approved	backend_reversed	denied	failed	processing	refunded	\
0	9.0	0.0	6.0	0.0	0.0	1.0	
1	13.0	0.0	8.0	0.0	0.0	0.0	
2	11.0	0.0	7.0	0.0	0.0	1.0	
3	12.0	0.0	3.0	0.0	0.0	0.0	
4	11.0	0.0	0.0	0.0	0.0	0.0	
...		
2636	31.0	0.0	8.0	0.0	0.0	0.0	
2637	34.0	0.0	8.0	0.0	0.0	0.0	
2638	40.0	0.0	11.0	0.0	0.0	0.0	
2639	35.0	0.0	4.0	0.0	0.0	0.0	
2640	38.0	0.0	10.0	0.0	0.0	0.0	

	reversed	total	denied_approved	denied_pct	...	hour_14	hour_15	\
0	0.0	16.0	66.659260	37.500000	...	0	0	
1	0.0	21.0	61.533728	38.095238	...	0	0	
2	0.0	19.0	63.630579	36.842105	...	0	0	
3	0.0	15.0	24.997917	20.000000	...	0	0	
4	0.0	11.0	0.000000	0.000000	...	0	0	
...			
2636	0.0	39.0	25.805619	20.512821	...	0	0	
2637	0.0	42.0	23.528720	19.047619	...	0	0	
2638	0.0	51.0	27.499313	21.568627	...	0	0	
2639	0.0	39.0	11.428245	10.256410	...	0	0	
2640	0.0	48.0	26.315097	20.833333	...	0	0	

	hour_16	hour_17	hour_18	hour_19	hour_20	hour_21	hour_22	hour_23
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0
...		
2636	0	0	0	0	0	0	0	1
2637	0	0	0	0	0	0	0	1
2638	0	0	0	0	0	0	0	1
2639	0	0	0	0	0	0	0	1
2640	0	0	0	0	0	0	0	1

[2641 rows x 87 columns]

```
[27]: transactions.columns
```

```
[27]: Index(['approved', 'backend_reversed', 'denied', 'failed', 'processing',
        'refunded', 'reversed', 'total', 'denied_approved', 'denied_pct',
        'reversed_approved', 'reversed_pct', 'failed_approved', 'failed_pct',
        'processing_approved', 'processing_pct', 'refunded_approved',
        'refunded_pct', 'backend_reversed_approved', 'backend_reversed_pct',
        'denied_sqr', 'reversed_sqr', 'failed_sqr', 'processing_sqr',
        'refunded_sqr', 'backend_reversed_sqr', 'denied_sqroot',
        'reversed_sqroot', 'failed_sqroot', 'processing_sqroot',
        'refunded_sqroot', 'backend_reversed_sqroot',
        'denied_approved_roll30_corr', 'reversed_approved_roll30_corr',
        'failed_approved_roll30_corr', 'processing_approved_roll30_corr',
        'refunded_approved_roll30_corr',
        'backend_reversed_approved_roll30_corr', 'denied_roll30_cumsum',
        'reversed_roll30_cumsum', 'failed_roll30_cumsum',
        'processing_roll30_cumsum', 'refunded_roll30_cumsum',
        'backend_reversed_roll30_cumsum', 'denied_roll30_avg',
        'reversed_roll30_avg', 'failed_roll30_avg', 'processing_roll30_avg',
        'refunded_roll30_avg', 'backend_reversed_roll30_avg',
        'denied_roll30_min', 'reversed_roll30_min', 'failed_roll30_min',
        'processing_roll30_min', 'refunded_roll30_min',
        'backend_reversed_roll30_min', 'denied_roll30_max',
        'reversed_roll30_max', 'failed_roll30_max', 'processing_roll30_max',
        'refunded_roll30_max', 'backend_reversed_roll30_max',
        'denied_sqr_distance', 'hour_00', 'hour_01', 'hour_02', 'hour_03',
        'hour_04', 'hour_05', 'hour_06', 'hour_07', 'hour_08', 'hour_09',
        'hour_10', 'hour_11', 'hour_12', 'hour_13', 'hour_14', 'hour_15',
        'hour_16', 'hour_17', 'hour_18', 'hour_19', 'hour_20', 'hour_21',
        'hour_22', 'hour_23'],
        dtype='object')
```

```
[28]: for col in transactions.columns:
        pct = transactions[f'{col}'].isnull().sum() * 100 / len(transactions.index)
        print(f'{col} - {pct}%')
```

```
approved - 0.0%
backend_reversed - 0.0%
denied - 0.0%
failed - 0.0%
processing - 0.0%
refunded - 0.0%
reversed - 0.0%
total - 0.0%
denied_approved - 0.0%
```

denied_pct - 0.0%
reversed_approved - 0.0%
reversed_pct - 0.0%
failed_approved - 0.0%
failed_pct - 0.0%
processing_approved - 0.0%
processing_pct - 0.0%
refunded_approved - 0.0%
refunded_pct - 0.0%
backend_reversed_approved - 0.0%
backend_reversed_pct - 0.0%
denied_sqr - 0.0%
reversed_sqr - 0.0%
failed_sqr - 0.0%
processing_sqr - 0.0%
refunded_sqr - 0.0%
backend_reversed_sqr - 0.0%
denied_sqroot - 0.0%
reversed_sqroot - 0.0%
failed_sqroot - 0.0%
processing_sqroot - 0.0%
refunded_sqroot - 0.0%
backend_reversed_sqroot - 0.0%
denied_approved_roll30_corr - 0.0%
reversed_approved_roll30_corr - 0.0%
failed_approved_roll30_corr - 0.0%
processing_approved_roll30_corr - 0.0%
refunded_approved_roll30_corr - 0.0%
backend_reversed_approved_roll30_corr - 0.0%
denied_roll30_cumsum - 0.0%
reversed_roll30_cumsum - 0.0%
failed_roll30_cumsum - 0.0%
processing_roll30_cumsum - 0.0%
refunded_roll30_cumsum - 0.0%
backend_reversed_roll30_cumsum - 0.0%
denied_roll30_avg - 0.0%
reversed_roll30_avg - 0.0%
failed_roll30_avg - 0.0%
processing_roll30_avg - 0.0%
refunded_roll30_avg - 0.0%
backend_reversed_roll30_avg - 0.0%
denied_roll30_min - 0.0%
reversed_roll30_min - 0.0%
failed_roll30_min - 0.0%
processing_roll30_min - 0.0%
refunded_roll30_min - 0.0%
backend_reversed_roll30_min - 0.0%
denied_roll30_max - 0.0%

```
reversed_roll30_max - 0.0%
failed_roll30_max - 0.0%
processing_roll30_max - 0.0%
refunded_roll30_max - 0.0%
backend_reversed_roll30_max - 0.0%
denied_sqr_distance - 0.0%
hour_00 - 0.0%
hour_01 - 0.0%
hour_02 - 0.0%
hour_03 - 0.0%
hour_04 - 0.0%
hour_05 - 0.0%
hour_06 - 0.0%
hour_07 - 0.0%
hour_08 - 0.0%
hour_09 - 0.0%
hour_10 - 0.0%
hour_11 - 0.0%
hour_12 - 0.0%
hour_13 - 0.0%
hour_14 - 0.0%
hour_15 - 0.0%
hour_16 - 0.0%
hour_17 - 0.0%
hour_18 - 0.0%
hour_19 - 0.0%
hour_20 - 0.0%
hour_21 - 0.0%
hour_22 - 0.0%
hour_23 - 0.0%
```

5 Principal Component Analysis (PCA)

We will use PCA to analyse the importance of the features we created and reduce them to a smaller number of features. We can also try to detect anomalies with PCA through **reconstruction error** when reconstructing the features after reducing them through PCA.

We will use PCA and keep the features that sum a 95% of variance explained

```
[29]: from sklearn.decomposition import PCA
      from sklearn.preprocessing import MinMaxScaler

      scaler = MinMaxScaler()

      x = scaler.fit_transform(transactions)

      pca = PCA(.95).set_output(transform='pandas')
```

```
reduced = pca.fit_transform(x)

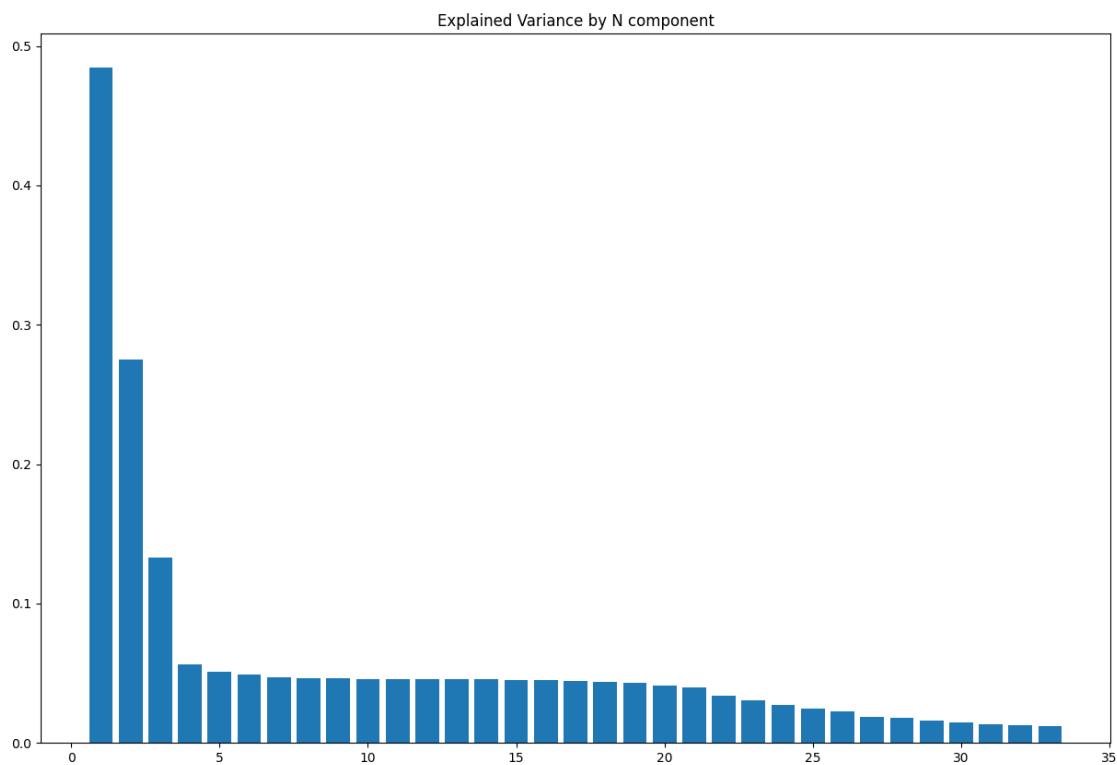
print('Features Before PCA: ',x.shape[1])
print('Features After PCA: ',reduced.shape[1])
```

Features Before PCA: 87

Features After PCA: 33

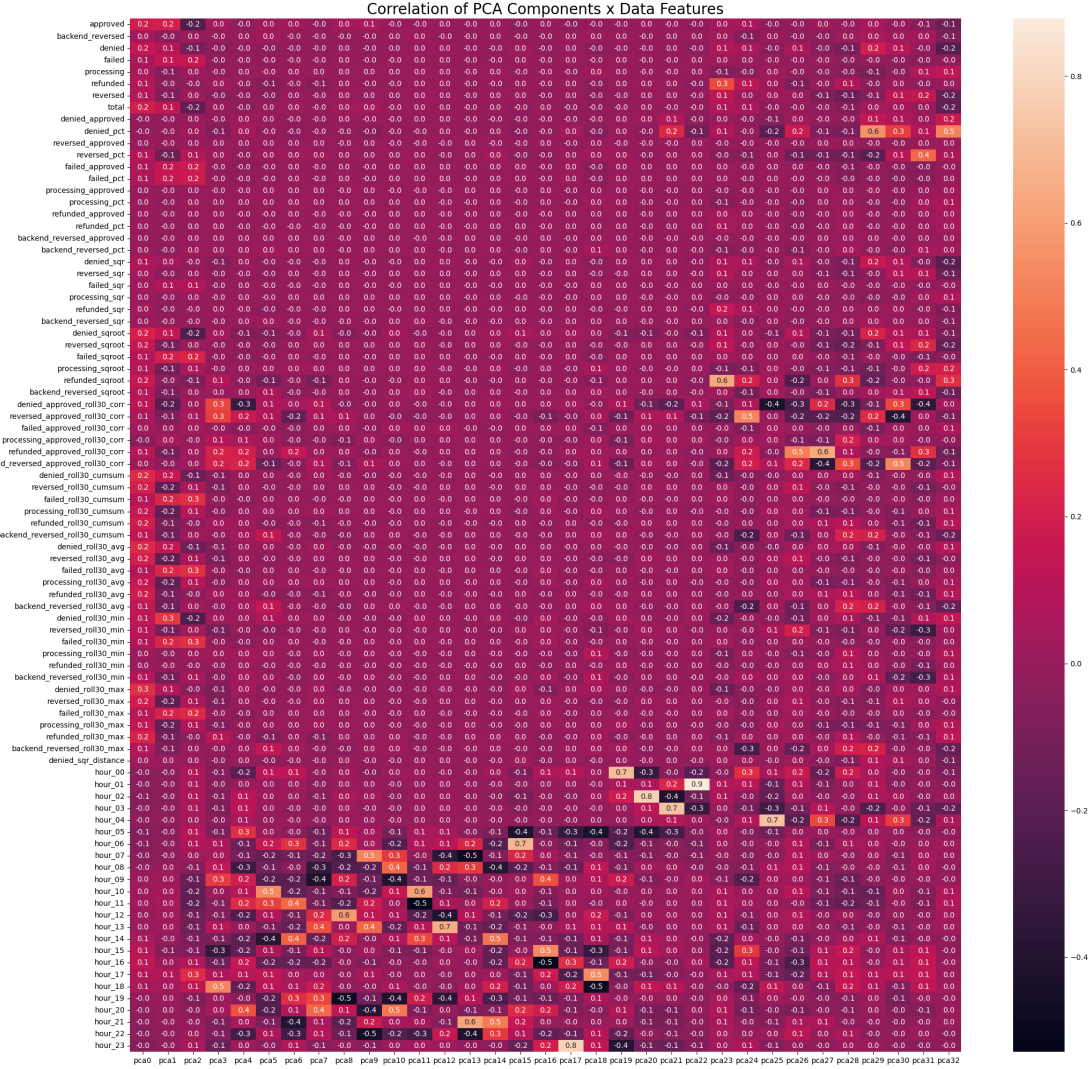
```
[30]: plt.figure(figsize=(15,10))
plt.bar(range(1, len(pca.explained_variance_) +1), pca.explained_variance_ )
plt.title('Explained Variance by N component')
```

```
[30]: Text(0.5, 1.0, 'Explained Variance by N component')
```



```
[31]: fig = plt.figure(figsize=(25,25))
components_df = pd.DataFrame(pca.components_ ,columns=transactions.columns,
    ↪ index =pca.get_feature_names_out())
sns.heatmap(components_df.T, xticklabels=True, yticklabels=True, annot=True,
    ↪fmt='.1f')
plt.title('Correlation of PCA Components x Data Features', fontsize=20)
```

```
[31]: Text(0.5, 1.0, 'Correlation of PCA Components x Data Features')
```




5.1 Reconstruction Error

Now that we reduced the data, we can reconstruct it through PCA and find errors in reconstruction. By subtracting the original scaled data from the reconstructed data we get the reconstruction error.

```
[32]: reverse_construct = pca.inverse_transform(reduced)
reconstruction_error = (reverse_construct - x)
reconstruction_error.columns = transactions.columns

plt.figure(figsize=(15,20))
plt.barh(reconstruction_error.columns, reconstruction_error.sum() )
plt.xticks(rotation=90)
plt.title('Sum of errors per Feature', fontsize=20)
```


[32]: Text(0.5, 1.0, 'Sum of errors per Feature')



```
[33]: plt.figure(figsize=(20,10))
fig = go.Figure()

for col in reconstruction_error.columns:
    fig.add_trace(go.Scatter(x=reconstruction_error.index,
        y=reconstruction_error[col],
        mode='lines',
```

```

        name=col))

fig.update_layout(
    autosize=False,
    width=1200,
    height=900,
)

fig.update_layout(
    title=dict(text="Reconstruction Error over time", font=dict(size=30))
)

fig.show()

```

<Figure size 2000x1000 with 0 Axes>

5.2 Looking Good

Our PCA model has some clear trends and anomalies detected, this is very promising for our anomaly detection and alert system. Let's check the average and the standard deviation of the errors so that maybe we can use the [3-sigma rule](#) to rule out the anomalies

Average of the averages

```
[34]: reconstruction_error.mean().mean()
```

```
[34]: np.float64(-6.168952615209917e-18)
```

Standard Deviations

```
[35]: pd.DataFrame(reconstruction_error.std())
```

```
[35]:
```

	0
approved	0.040513
backend_reversed	0.034633
denied	0.032984
failed	0.026389
processing	0.036593
...	...
hour_19	0.008720
hour_20	0.006816
hour_21	0.005729
hour_22	0.006516
hour_23	0.007740

[87 rows x 1 columns]

Average of Standard Deviations

```
[36]: reconstruction_error.std().mean()
```

```
[36]: np.float64(0.02927367747523001)
```

We will square the errors, so we can further separate normal data from anomalies, this will also make all values positive. Then we take the average and the mean and use the 3-sigma rule to plot the anomalies

```
[37]: re_squared = reconstruction_error**2

def rule_out(reconstruction_error:pd.Series, re_squared:pd.DataFrame):
    column = reconstruction_error.name
    mean = re_squared[column].mean()
    std = re_squared[column].std()

    sigma3 = mean + 3*std
    sigma3_negative = mean - 3*std

    return (reconstruction_error >= sigma3) | (reconstruction_error <=
↪sigma3_negative)
```

```
[38]: plt.figure(figsize=(20,10))
fig = go.Figure()

for col in re_squared.columns:
    anomalies = re_squared[col].mask( ~rule_out(re_squared[col], re_squared) )
    fig.add_trace(go.Scatter(x=reconstruction_error.index, y=anomalies,
                             mode='markers',
                             name=col))

fig.update_layout(
    autosize=False,
    width=1200,
    height=900,
)

fig.update_layout(
    title=dict(text="Reconstruction Error over time", font=dict(size=30))
)

fig.show()
```

<Figure size 2000x1000 with 0 Axes>

5.3 Disturbance Readings

We will use **Softmax** to see how much each feature contributes to the total of the reconstruction error, then we will stack the values to create a kind of *seismometer*. Unfortunately we can't put the labels because there are too many of them and it wouldn't help to understand the data, but you will see that it does help to identify anomalous behavior

```
[39]: from scipy.special import softmax
plt.figure(figsize=(20,10))

labels = re_squared.columns

re_squared = pd.DataFrame(softmax(re_squared, axis=1), columns=labels)

plt.stackplot(re_squared.index, *[re_squared[column] for column in labels],
              labels = labels)

plt.ylim(0,1)
```

[39]: (0.0, 1.0)




The small *scribbles* you see are from anomalies that bump the other stacks because of it's higher contribution towards the total error. As softmax adds everything up to 1, it's natural that the other features lose space in the readings. That's why this graph helps a lot in finding anomalies

```
[40]: plt.figure(figsize=(20,10))
plt.stackplot(re_squared.index, *[re_squared[column] for column in labels],
              labels = labels)

plt.xlim(800, 1000)
plt.ylim(0,1)
```

[40]: (0.0, 1.0)



It might seem small, but remember these are 87 features. With softmax making all the feature errors summing to 1, that leaves about ~1% for each feature to contribute

```
[41]: re_squared.mean()
```

```
[41]: approved          0.011500
      backend_reversed  0.011496
      denied            0.011494
      failed            0.011489
      processing        0.011497
      ...
      hour_19           0.011482
      hour_20           0.011482
      hour_21           0.011482
      hour_22           0.011482
      hour_23           0.011482
      Length: 87, dtype: float64
```

```
[42]: re_squared.max().sort_values(ascending=False)
```

```
[42]: processing_approved      0.028251
      refunded_rol130_min     0.027583
      backend_reversed_approved 0.027390
      reversed_approved       0.027317
      refunded_approved       0.027189
      ...
      hour_19                 0.011494
      hour_08                 0.011494
      hour_20                 0.011494
      hour_21                 0.011494
      hour_22                 0.011494
      Length: 87, dtype: float64
```

As we can see, the average of the error values is around 1.14% which is very small. But the max values i.e. the main anomalous values go up to almost 3% which is a large difference from the normal readings

5.4 Why PCA

The reason we will stick with using PCA for this and not using other alternative models for anomaly detection is the following:

- PCA can handle these smaller datasets, which other models tend to not perform well
- Our PCA approach has a very good “lead” on spotting anomalies
- There is a hand few of other anomaly detection algorithms that could work, but testing each approach would require time we don’t have

6 Final Touches

We will do the final touches so we can focus on the API and systems. Creating functions to receive the data, transform and detect anomalies. Saving the models to use later, etc.

Dict to save means and standard deviations for anomaly detection

```
[43]: re_squared = reconstruction_error**2

means_dict = {}
for col in re_squared.columns:
    mean = re_squared[col].mean()
    std = re_squared[col].std()
    means_dict.update(
        {col: {'mean': mean, 'std': std}}
    )
```

Function to receive raw data in a API and detect anomalies, returning a dictionary

```
[44]: def detect_anomalies(data:pd.DataFrame):
    anomaly_dict = {}

    columns = data.columns
    data = scaler.transform(data)
    pca_transformed = pca.transform(data)
    pca_inversed = pca.inverse_transform(pca_transformed)
    pca_inversed.columns = data.columns
    data = (data - pca_inversed)**2
    softmax_data = softmax(data)

    for col in data.columns:
        mean = means_dict.get(col).get('mean')
        std = means_dict.get(col).get('std')

        sigma3 = mean + 3*std
        sigma3_negative = mean - 3*std

        value = data[col].values[0]
        key = col.split('_')[0]
        pct = softmax_data[col].values[0]

        if value >= sigma3 or value <= sigma3_negative: anomaly = True

        else: anomaly = False

        anomaly_dict.update({key: {'anomaly':anomaly, 'value':value,'pct':pct}})

    return anomaly_dict
```

Saving models and important data in files

```
[45]: import joblib
import pickle

joblib.dump(scaler, 'scaler.save')
joblib.dump(pca, 'pca.save')
joblib.dump(linearR_denied, 'linearR_denied.save')
```

```
with open('means_dict.pkl', 'wb') as f:  
    pickle.dump(means_dict, f)
```